


# TIME-WARPING WITH CATS EFFECT

Jacob Wang

# Hello

- Software Developer at  medidata
- @jatcwang
- Talk to me about types, FP, libraries and tools :)

**It's Tuesday afternoon...**

# It's Tuesday afternoon...

- You just finished the beautiful data-processing pipeline using *fs2*, with windowed batched message acknowledgements, exponential backoff on error, etc etc.

# It's Tuesday afternoon...

- You just finished the beautiful data-processing pipeline using *fs2*, with windowed batched message acknowledgements, exponential backoff on error, etc etc.
- But how do you test it?

# The Problem

## Tests involving time and delays are difficult to write

- Using `Thread.sleep` (or any kind of real clock) leads to long-running and flaky tests
- but what else can we do?

# In this talk

- Threads, thread pools and schedulers from the ground up
- How to build a “controlled” thread pool / scheduler
- How to write tests with cats-effect’s TestContext

# Back to the basics

In the beginning, we have `java.lang.Thread`

```
trait Runnable {  
  def run(): Unit  
}  
  
val t = new Thread(new Runnable {  
  def run() = {  
    Thread.sleep(1000)  
    println("hi!") }  
})  
t.run() // "hi!" after a delay of ~1 second
```



# Back to the basics

In the beginning, we have `java.lang.Thread`

- `new Thread(...)` creates a new JVM thread, backed by a native OS thread

```
trait Runnable {  
  def run(): Unit  
}  
  
val t = new Thread(new Runnable {  
  def run() = {  
    Thread.sleep(1000)  
    println("hi!") }  
})  
t.run() // "hi!" after a delay of ~1 second
```

# Back to the basics

In the beginning, we have `java.lang.Thread`

- `new Thread(...)` creates a new JVM thread, backed by a native OS thread
  - `thread.run()` executes the Runnable initially assigned to the thread

```
trait Runnable {  
  def run(): Unit  
}  
  
val t = new Thread(new Runnable {  
  def run() = {  
    Thread.sleep(1000)  
    println("hi!") }  
})  
t.run() // "hi!" after a delay of ~1 second
```

# Back to the basics

In the beginning, we have `java.lang.Thread`

- `new Thread(...)` creates a new JVM thread, backed by a native OS thread
  - `thread.run()` executes the Runnable initially assigned to the thread
- `Thread.sleep()` suspends the current thread's execution, and OS schedules other threads to run

```
trait Runnable {  
  def run(): Unit  
}  
  
val t = new Thread(new Runnable {  
  def run() = {  
    Thread.sleep(1000)  
    println("hi!")  
  }  
})  
t.run() // "hi!" after a delay of ~1 second
```

# What's wrong with Threads?

# What's wrong with Threads?

- Clunky API

# What's wrong with Threads?

- Clunky API
- Overhead both in processing time and memory

# What's wrong with Threads?

- Clunky API
- Overhead both in processing time and memory
  - Memory overhead (each thread has its own stack)

# What's wrong with Threads?

- Clunky API
- Overhead both in processing time and memory
  - Memory overhead (each thread has its own stack)
  - “Context Switching” cost from one thread to another



# What's wrong with Threads?

- Clunky API
- Overhead both in processing time and memory
  - Memory overhead (each thread has its own stack)
  - “Context Switching” cost from one thread to another
  - In many use cases it limits scalability  
(e.g. Backend services serving a lot of requests all hitting the DB)

# Thread Pools (ExecutorService)

```
trait ExecutorService {  
  def submit(runnable: Runnable): Unit  
}
```

# Thread Pools (ExecutorService)

```
trait ExecutorService {  
  def submit(runnable: Runnable): Unit  
}
```

- A few **worker threads** take from a queue of **Runnable**s and execute them

# Thread Pools (ExecutorService)


```
trait ExecutorService {  
  def submit(runnable: Runnable): Unit  
}
```

- A few **worker threads** take from a queue of **Runnable**s and execute them
- **ExecutionContext** is Scala's own simpler interface for the same idea

# Thread Pools (ExecutorService)

```
trait ExecutorService {  
  def submit(runnable: Runnable): Unit  
}
```

- A few **worker threads** take from a queue of **Runnables** and execute them
- **ExecutionContext** is Scala's own simpler interface for the same idea



# ScheduledExecutorService

```
trait ScheduledExecutorService extends ExecutorService {  
  // Simplified, delay in nanoseconds  
  def schedule(runnable: Runnable, delay: Long): Unit  
}
```

# ScheduledExecutorService

```
trait ScheduledExecutorService extends ExecutorService {  
  // Simplified, delay in nanoseconds  
  def schedule(runnable: Runnable, delay: Long): Unit  
}
```

- Schedule a **Runnable** to be executed after a delay

# ScheduledExecutorService

```
trait ScheduledExecutorService extends ExecutorService {  
  // Simplified, delay in nanoseconds  
  def schedule(runnable: Runnable, delay: Long): Unit  
}
```

- Schedule a **Runnable** to be executed after a delay
- Typically backed by a single max-priority thread running (essentially) `Thread.sleep` continuously



# ScheduledExecutorService

```
trait ScheduledExecutorService extends ExecutorService {  
  // Simplified, delay in nanoseconds  
  def schedule(runnable: Runnable, delay: Long): Unit  
}
```

- Schedule a **Runnable** to be executed after a delay
- Typically backed by a single max-priority thread running (essentially) `Thread.sleep` continuously
  - See `java.util.concurrent.ScheduledThreadPoolExecutor`

# One Simple Idea

Can we write a `ScheduledExecutorService` that doesn't use sleep nor run anything, but instead just record the `Runnable`s and when it should be run?

# One Simple Idea

Can we write a `ScheduledExecutorService` that doesn't use sleep nor run anything, but instead just record the `Runnable`s and when it should be run?

```
final case class Task(runnable: Runnable, runAt: Long)

class ControlledScheduledExecutorService extends ScheduledExecutorService {
  private var currentTime: Long = 0
  private var allTasks: SortedSet[Task] = SortedSet.empty // sorted by runAt

  override def schedule(runnable: Runnable, delay: Long): Unit = {
    allTasks = allTasks :+ Task(runnable, currentTime + delay)
  }
}
```

# Preparing to time-warp...

Since no OS thread nor scheduling are involved, we have total control over what to run and when to run it

```
class ControlledScheduledExecutorService extends ScheduledExecutorService {  
  private var currentTime: Long = 0  
  private var allTasks: SortedSet[Task] = SortedSet.empty // sorted by runAt  
  
  override def schedule(runnable: Runnable, delay: Long): Unit = // ...  
  
  def tick(elapseTime: Long) = {  
    currentTime = currentTime + elapseTime  
    var nextTask = allTasks.findAndRemove(_.runAt <= currentTime)  
    while (nextTask.nonEmpty) {  
      nextTask.get.runnable.run()  
      nextTask = allTasks.findAndRemove(_.runAt <= currentTime)  
    }  
  }  
}
```

# Time-warp complete!

And that's all there is to it!

# Time-warp complete!

And that's all there is to it!

- Evaluation depends on “what task is in the queue, and when should they be run”, which means there's no difference between running 100 **Runnable** scheduled all within 1 millis or 1000 years!

# Time-warp complete!

And that's all there is to it!

- Evaluation depends on “what task is in the queue, and when should they be run”, which means there's no difference between running 100 **Runnable** scheduled all within 1 millis or 1000 years!
- “Great, but only cavemen use Runnables”

# It's Runnables all the way down

Almost all effect libraries / Futures ultimately ends up as **Runnables** being executed un by a thread pool.

```
implicit val ec: ExecutionContext = ...
Future {
  1 + 1 // 1
}.map { res =>
  println(res) // 2
}

// Eventually translates to something like...
ec.execute(new Runnable {
  def run() = {
    val res = 1 + 1 // 1
    ec.execute(new Runnable { // submit the "next step" back to the EC
      def run() = {
        println(res) // 2
      }
    })
  }
})
})
```



# It's Runnables all the way down

Almost all effect libraries / Futures ultimately ends up as **Runnables** being executed un by a thread pool.

```
implicit val ec: ExecutionContext = ...
Future {
  1 + 1 // 1
}.map { res =>
  println(res) // 2
}

// Eventually translates to something like...
ec.execute(new Runnable {
  def run() = {
    val res = 1 + 1 // 1
    ec.execute(new Runnable { // submit the "next step" back to the EC
      def run() = {
        println(res) // 2
      }
    })
  }
})
})
```

- So the technique covered thus far can be applied everywhere!

# cats.effect TestContext

**TestContext** is one implementation of controlled thread pool / scheduler

```
val testCtx = cats.effect.laws.util.TestContext() // TestContext <: ExecutionContext
implicit val ctxShift: ContextShift[IO] = testCtx.ioContextShift
implicit val timer: Timer[IO] = testCtx.ioTimer

val io = for {
  _ <- IO { println("hi!") }
  _ <- IO.sleep(100.days)
  _ <- IO { println("I have awakened!") }
}

// execute our IO in the TestContext
ctxShift.evalOn(testCtx)(io).unsafeRunAsyncAndForget() // Step 2 "run" the IO

println("Ticking 50 days..") // Step 3 time-warp!
testCtx.tick(50.days) // "hi!" printed here

println("Ticking another 50 days..")
testCtx.tick(50.days) // "I have awakened!" is printed
```

# cats.effect TestContext

**TestContext** is one implementation of controlled thread pool / scheduler

- You can find it in **cats-effect-laws** library dependency

```
val testCtx = cats.effect.laws.util.TestContext() // TestContext <: ExecutionContext
implicit val ctxShift: ContextShift[IO] = testCtx.ioContextShift
implicit val timer: Timer[IO] = testCtx.ioTimer

val io = for {
  _ <- IO { println("hi!") }
  _ <- IO.sleep(100.days)
  _ <- IO { println("I have awakened!") }
}

// execute our IO in the TestContext
ctxShift.evalOn(testCtx)(io).unsafeRunAsyncAndForget() // Step 2 "run" the IO

println("Ticking 50 days..") // Step 3 time-warp!
testCtx.tick(50.days) // "hi!" printed here

println("Ticking another 50 days..")
testCtx.tick(50.days) // "I have awakened!" is printed
```

# Structuring your test

How to structure our time-sensitive tests?

# Structuring your test

How to structure our time-sensitive tests?

- We have “app logic” and “assertions” (observer)

# Structuring your test

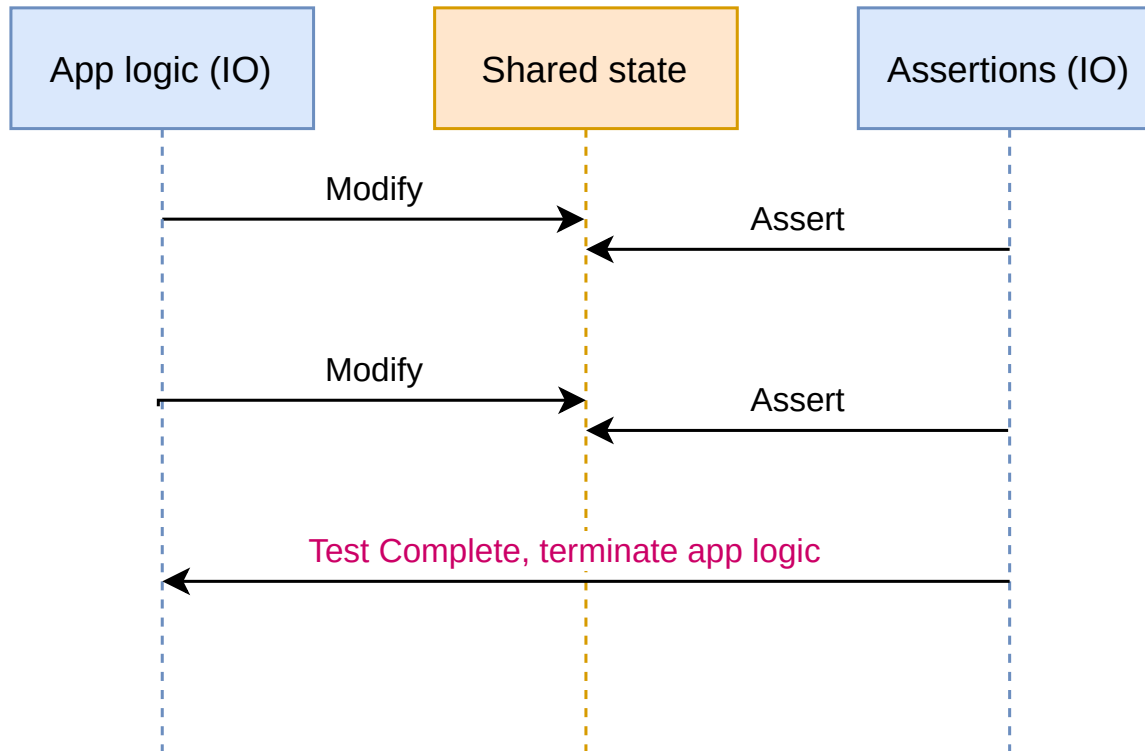
How to structure our time-sensitive tests?

- We have “app logic” and “assertions” (observer)
- Run these two IOs in “parallel”

# Structuring your test

How to structure our time-sensitive tests?

- We have “app logic” and “assertions” (observer)
- Run these two IOs in “parallel”



```
val appLogic: IO[Unit] = ???           // Step 1: Define Logic and Assertions
val assertions: IO[Unit] = for {
  _ <- IO.sleep(5.seconds)
  _ <- assertState(state1)
  _ <- IO.sleep(5.seconds)
  _ <- assertState(state2)
} yield succeed

// Step 2: Run it with test "framework" (reusable across test suites)
val testContext = TestContext()
val test = for {
  fibre <- appLogic.start // forks the app logic into the background
  _ <- IO.sleep(1.nanos) // very small delay to ensure assertions
                        // always run after app logic has happened
  _ <- assertions.guarantee(fibre.cancel)
} yield succeed

// Step 3: "Execute" the test
val f: Future[Assertion] = test.unsafeToFuture()

// Actually execute all the app logic and assertions
// Note that assertions will terminate the appLogic when it is done
// so that this works for infinitely repeating app logic too
testContext.tick(1000.days)
```



# Final words

# Final words

- The underlying works is based on JVM primitives , you can time-warp with Scala Future, Monix, ZIO, and Java `CompletableFuture` too ;)

# Final words

- The underlying works is based on JVM primitives , you can time-warp with Scala Future, Monix, ZIO, and Java `CompletableFuture` too ;)
- Since only TestContext is controlled, **ALL** your Runnables need to be submitted to it

# Final words

- The underlying works is based on JVM primitives , you can time-warp with Scala Future, Monix, ZIO, and Java `CompletableFuture` too ;)
- Since only TestContext is controlled, **ALL** your Runnables need to be submitted to it
  - Doobie Transactor, any `Future{..}/future.map`, etc etc

# Final words

- The underlying works is based on JVM primitives , you can time-warp with Scala Future, Monix, ZIO, and Java `CompletableFuture` too ;)
- Since only TestContext is controlled, **ALL** your Runnables need to be submitted to it
  - Doobie Transactor, any `Future{..}/future.map`, etc etc
  - External libraries should allow you to specify a thread pool

# Acknowledgements

- People on cats-effect gitter chat who helped me setup TestContext tests
  - Daniel Spiewak (@djspiwak)
  - Fabio Labella (aka SystemFw)
  - Gavin Bisesi (@Daenyth)

**Thank you!**

// reveal.js plugins