

Deep dive into Context Propagation


Jacob Wang

London Scala User Group, 2024 March

This talk


- Mechanisms for context propagation
- **otel4s** and **OpenTelemetry**

Hello


- Work at  medidata
- Maintainer of Doobie and author of Difflicious + other libs
- <https://mas.to/@jatcwang>

Context Propagation - Large and Small

Context Propagation - Large and Small




Context Propagation - Large and Small




- Context: Not used for fulfilling the request, but instead for monitoring and debugging

Context Propagation - Large and Small




- Context: Not used for fulfilling the request, but instead for monitoring and debugging
 - e.g. Trace ID, Tenant ID, User ID

Context Propagation - Large and Small



- Context: Not used for fulfilling the request, but instead for monitoring and debugging
 - e.g. Trace ID, Tenant ID, User ID
- This talk: Passing context *within* an application

Context Propagation - Large and Small



- Context: Not used for fulfilling the request, but instead for monitoring and debugging
 - e.g. Trace ID, Tenant ID, User ID
- This talk: Passing context *within* an application
- Invisible / Non-local mechanisms only

java.lang.ThreadLocal

java.lang.ThreadLocal

- Store state **local** to the code executing on that thread only

java.lang.ThreadLocal

- Store state **local** to the code executing on that thread only
- ThreadLocal instances are “key” to set/remove their corresponding value stored in the thread

java.lang.ThreadLocal

- Store state **local** to the code executing on that thread only
- ThreadLocal instances are “key” to set/remove their corresponding value stored in the thread
- Conceptually, each Thread has a `Map[ThreadLocal[A], A]`

java.lang.ThreadLocal

- Store state **local** to the code executing on that thread only
- ThreadLocal instances are “key” to set/remove their corresponding value stored in the thread
- Conceptually, each Thread has a `Map[ThreadLocal[A], A]`

```
1  val TL_F00: ThreadLocal[Int] = ThreadLocal.withInitial(() => 0)
2  val TL_BAR: ThreadLocal[Int] = ThreadLocal.withInitial(() => 0)
3
4  TL_F00.set(5)
5
6  TL_F00.get() // == 5
7  TL_BAR.get() // == 0, because TL_BAR is a different "key" from TL_F00
8
9  TL_F00.remove()
10
11 TL_F00.get() // == 0
```

```
1  val t1 = new Thread(() => {
2      TL_F00.get()    // == 0
3      TL_F00.set(1)
4      TL_F00.get()    // == 1
5  }, "t1").start()
6
7  val t2 = new Thread(() => {
8      TL_F00.get()    // == 0
9      TL_F00.set(3)
10     TL_F00.get()    // == 3
11 }, "t2").start()
```


java.lang.ThreadLocal

- Allow us to pass context without explicit parameters

java.lang.ThreadLocal

- Allow us to pass context without explicit parameters
- Java logging **Mapped Diagnostic Context** (MDC) and **OpenTelemetry** pass their context using ThreadLocal

java.lang.ThreadLocal

- Allow us to pass context without explicit parameters
- Java logging **Mapped Diagnostic Context** (MDC) and **OpenTelemetry** pass their context using ThreadLocal
- Can use `InheritableThreadLocal` to pass on context to child threads

java.lang.ThreadLocal

- Allow us to pass context without explicit parameters
- Java logging **Mapped Diagnostic Context** (MDC) and **OpenTelemetry** pass their context using ThreadLocal
- Can use `InheritableThreadLocal` to pass on context to child threads
 - e.g. background tasks

The trouble with (many) Threads

The trouble with (many) Threads

- Old-school HTTP libraries create a new thread per request

The trouble with (many) Threads

- Old-school HTTP libraries create a new thread per request
- Each thread has a base cost of ~1MB

The trouble with (many) Threads

- Old-school HTTP libraries create a new thread per request
- Each thread has a base cost of ~1MB
- Threads are often blocked

The trouble with (many) Threads

- Old-school HTTP libraries create a new thread per request
- Each thread has a base cost of ~1MB
- Threads are often blocked
- CPU cores switching between threads (“context switching”) are expensive

Let's reuse threads! 💡

Let's reuse threads! 💡

- **threadpool**: pool of reusable threads and submit “tasks” (`Runnable`) to it

Let's reuse threads! 💡

- **threadpool**: pool of reusable threads and submit “tasks” (`Runnable`) to it
- `Runnable` (equivalent to `() => Unit`) have much smaller overhead

Let's reuse threads! 💡

- **threadpool**: pool of reusable threads and submit “tasks” (`Runnable`) to it
- `Runnable` (equivalent to `() => Unit`) have much smaller overhead
- “Async”

Let's reuse threads! 💡

- **threadpool**: pool of reusable threads and submit “tasks” (`Runnable`) to it
- `Runnable` (equivalent to `() => Unit`) have much smaller overhead
- “Async”
- Java/Scala Futures and Effect runtimes (Cats-effect, ZIO) are all built on top of threadpools

Let's reuse threads! 💡

- **threadpool**: pool of reusable threads and submit “tasks” (`Runnable`) to it
- `Runnable` (equivalent to `() => Unit`) have much smaller overhead
- “Async”
- Java/Scala Futures and Effect runtimes (Cats-effect, ZIO) are all built on top of threadpools

```
1  trait Runnable {  
2    def run(): Unit  
3  }  
4  
5  trait Executor { // The threadpool interface  
6    def execute(task: Runnable): Unit  
7  }
```

```
1  val threadpool = Executors.newFixedThreadPool(4)
2  val httpClient = ...
3  val finishCallback: Result => Unit = ...
4
5  threadpool.execute(() => {
6      println(s"step 1")
7      val httpRequest = ...
8
9      httpClient.send(httpRequest, onResponse = response => {
10         // Callback will submit next task to threadpool
11         threadpool.execute(() => {
12             val result = doWork(response)
13             finishCallback(result)
14         })
15     })
16 })
17 }
```




I think I forgot something.



If you forgot, then
it wasn't important.



Yeah, you're right.

imgflip.com



ThreadLocal

Losing Context

```
1  val threadpool = Executors.newFixedThreadPool(4) // 4 threads in pool
2  val CONTEXT = ThreadLocal.withInitial(() => "no user")
3
4  threadpool.execute(() => {
5      CONTEXT.set("user1")
6      println(s"${Thread.currentThread().getName}: step 1 for $
XT.get()})")
7
8      threadpool.execute(() => {
9          println(s"${Thread.currentThread().getName}: step 2 for $
XT.get()})")
10
11          threadpool.execute(() => {
12              println(s"${Thread.currentThread().getName}: complete for $
XT.get()})")
13          })
14
15      })
16  })
```

Losing Context

```
1  val threadpool = Executors.newFixedThreadPool(4) // 4 threads in pool
2  val CONTEXT = ThreadLocal.withInitial(() => "no user")
3
4  threadpool.execute(() => {
5      CONTEXT.set("user1")
6      println(s"${Thread.currentThread().getName}: step 1 for $
XT.get()})")
7
8      threadpool.execute(() => {
9          println(s"${Thread.currentThread().getName}: step 2 for $
XT.get()})")
10
11          threadpool.execute(() => {
12              println(s"${Thread.currentThread().getName}: complete for $
XT.get()})")
13          })
14
15      })
16  })
```

```
pool-1-thread-1: step 1 for user1
pool-1-thread-2: step 2 for no_user
pool-1-thread-3: complete for no_user
```

Losing Context

```
1  val threadpool = Executors.newFixedThreadPool(4) // 4 threads in pool
2  val CONTEXT = ThreadLocal.withInitial(() => "no user")
3
4  threadpool.execute(() => {
5      CONTEXT.set("user1")
6      println(s"${Thread.currentThread().getName}: step 1 for $
XT.get()})")
7
8      threadpool.execute(() => {
9          println(s"${Thread.currentThread().getName}: step 2 for $
XT.get()})")
10
11          threadpool.execute(() => {
12              println(s"${Thread.currentThread().getName}: complete for $
XT.get()})")
13          })
14
15      })
16  })
```

```
pool-1-thread-1: step 1 for user1
pool-1-thread-2: step 2 for no_user
pool-1-thread-3: complete for no_user
```

- Lost the context because **another thread** executed step 2


Losing Context

```
1  val threadpool = Executors.newFixedThreadPool(4) // 4 threads in pool
2  val CONTEXT = ThreadLocal.withInitial(() => "no user")
3
4  threadpool.execute(() => {
5      CONTEXT.set("user1")
6      println(s"${Thread.currentThread().getName}: step 1 for $
XT.get()})")
7
8      threadpool.execute(() => {
9          println(s"${Thread.currentThread().getName}: step 2 for $
XT.get()})")
10
11          threadpool.execute(() => {
12              println(s"${Thread.currentThread().getName}: complete for $
XT.get()})")
13          })
14
15      })
16  })
```

```
pool-1-thread-1: step 1 for user1
pool-1-thread-2: step 2 for no_user
pool-1-thread-3: complete for no_user
```

- Lost the context because **another thread** executed step 2
- Worse, another unrelated task executing on thread 1 now has “user 1” as its context!

Making ThreadLocal work with threadpools

-  Capture ThreadLocal value in the submitting thread and set it when running in the new thread

Making ThreadLocal work with threadpools

- 💡 Capture ThreadLocal value in the submitting thread and set it when running in the new thread

```
1  class CurrentContextExecutor(delegate: Executor) extends Executor:
2    override def execute(task: Runnable): Unit =
3      val toAttach = CONTEXT.get()           // T1
4      val wrappedTask = wrap(task, toAttach) // T1
5      delegate.execute(wrappedTask)         // T1
6
7    def wrap(task: Runnable, toAttach: Context): Runnable =
8      () =>
9        try                                  // T2
10          CONTEXT.set(toAttach)             // T2
11          task.run()                        // T2
12        finally                             // T2
13          CONTEXT.remove()                  // T2
```


Making ThreadLocal work with threadpools

- 💡 Capture ThreadLocal value in the submitting thread and set it when running in the new thread

```
1  class CurrentContextExecutor(delegate: Executor) extends Executor:
2    override def execute(task: Runnable): Unit =
3      val toAttach = CONTEXT.get()           // T1
4      val wrappedTask = wrap(task, toAttach) // T1
5      delegate.execute(wrappedTask)         // T1
6
7    def wrap(task: Runnable, toAttach: Context): Runnable =
8      () =>
9        try                                  // T2
10          CONTEXT.set(toAttach)             // T2
11          task.run()                        // T2
12        finally                             // T2
13          CONTEXT.remove()                  // T2
```

- Resurface the context we want to pass on (Line 3)

Making ThreadLocal work with threadpools

- 💡 Capture ThreadLocal value in the submitting thread and set it when running in the new thread

```
1  class CurrentContextExecutor(delegate: Executor) extends Executor:
2    override def execute(task: Runnable): Unit =
3      val toAttach = CONTEXT.get()           // T1
4      val wrappedTask = wrap(task, toAttach) // T1
5      delegate.execute(wrappedTask)         // T1
6
7    def wrap(task: Runnable, toAttach: Context): Runnable =
8      () =>
9        try                                  // T2
10          CONTEXT.set(toAttach)             // T2
11          task.run()                        // T2
12        finally                             // T2
13          CONTEXT.remove()                 // T2
```

- Resurface the context we want to pass on (Line 3)
- Set the context we want to pass on (Line 10)

Making ThreadLocal work with threadpools

- 💡 Capture ThreadLocal value in the submitting thread and set it when running in the new thread

```
1  class CurrentContextExecutor(delegate: Executor) extends Executor:
2    override def execute(task: Runnable): Unit =
3      val toAttach = CONTEXT.get()           // T1
4      val wrappedTask = wrap(task, toAttach) // T1
5      delegate.execute(wrappedTask)         // T1
6
7    def wrap(task: Runnable, toAttach: Context): Runnable =
8      () =>
9        try                                  // T2
10          CONTEXT.set(toAttach)             // T2
11          task.run()                        // T2
12        finally                             // T2
13          CONTEXT.remove()                  // T2
```

- Resurface the context we want to pass on (Line 3)
- Set the context we want to pass on (Line 10)
- Remove the ThreadLocal value (Line 13)

cats.effect.IOLocal

cats.effect.IOLocal

- With Cats-Effect, *can* use `ThreadLocal` + task wrapping trick we just learned

cats.effect.IOLocal

- With Cats-Effect, *can* use `ThreadLocal` + task wrapping trick we just learned
- But we have a nicer solution: `IOLocal`

cats.effect.IOLocal

- With Cats-Effect, *can* use `ThreadLocal` + task wrapping trick we just learned
- But we have a nicer solution: `IOLocal`
- `IOLocal` allow us to pass context through the executing fiber

cats.effect.IOLocal

- With Cats-Effect, *can* use `ThreadLocal` + task wrapping trick we just learned
- But we have a nicer solution: `IOLocal`
- `IOLocal` allow us to pass context through the executing fiber
- On fork, child fibers inherit a copy of the parent's context

cats.effect.IOLocal Example

```
1  for
2    // Create the IOLocal "key"
3    CONTEXT <- IOLocal(default = 0)
4
5    _ <- CONTEXT.get           // == 0
6    _ <- CONTEXT.set(5)
7
8    // fork!
9    fiber1 <- (for {
10      _ <- CONTEXT.get       // == 5
11      _ <- CONTEXT.set(6)
12      _ <- CONTEXT.get       // == 6
13    } yield ()).start
14
15    _ <- CONTEXT.get           // == 5
16    _ <- CONTEXT.set(10)
17    _ <- fiber1.joinWithNever
18    _ <- CONTEXT.get           // == 10
19  yield ()
```

What we've seen so far

What we've seen so far

- ThreadLocal

What we've seen so far

- ThreadLocal
- Thread pools and how ThreadLocal can work with them

What we've seen so far

- ThreadLocal
- Thread pools and how ThreadLocal can work with them
- IOLocal for cats-effect

What we've seen so far

- ThreadLocal
- Thread pools and how ThreadLocal can work with them
- IOLocal for cats-effect
- Now, let's apply what we've learned

OpenTelemetry - Quick Intro

- **OpenTelemetry**: An open standard for Distributed Tracing, Metrics, etc

OpenTelemetry - Quick Intro

- **OpenTelemetry**: An open standard for Distributed Tracing, Metrics, etc
- **Span**: A recorded unit of work


OpenTelemetry - Quick Intro

- **OpenTelemetry:** An open standard for Distributed Tracing, Metrics, etc
- **Span:** A recorded unit of work
 - Attributes include `traceId`, `parentSpanId`, `isError` and `duration`


OpenTelemetry - Quick Intro

- **OpenTelemetry:** An open standard for Distributed Tracing, Metrics, etc
- **Span:** A recorded unit of work
 - Attributes include `traceId`, `parentSpanId`, `isError` and `duration`
- **Trace:** Links together a set of spans (same `traceId`), so we can track the execution trace from start to finish across services.

Trace Start **December 16, 2018 5:19 PM** Duration **700.68ms** Services **6** Depth **5** Total Spans **50**



Service & Operation	▼ > ⌵ ⌶	0ms	175.17ms	350.34ms	525.51ms	700.68ms
---------------------	---------	-----	----------	----------	----------	----------



Otel4s - OpenTelemetry for Scala

Otel4s - OpenTelemetry for Scala

- Typelevel ecosystem

Otel4s - OpenTelemetry for Scala

- Typelevel ecosystem
- Follows OTel terminology and specification, but does not directly use OpenTelemetry-Java types

Otel4s - OpenTelemetry for Scala

- Typelevel ecosystem
- Follows OTel terminology and specification, but does not directly use OpenTelemetry-Java types
- Uses OpenTelemetry-Java as backend (e.g. reporting spans)

OpenTelemetry - Instrumentation API concepts

OpenTelemetry - Instrumentation API concepts

- `Context`: Immutable key-value pairs for storing all OTel-related objects

OpenTelemetry - Instrumentation API concepts

- `Context`: Immutable key-value pairs for storing all OTel-related objects
 - The `Context` object itself is propagated around (e.g. `ThreadLocal`, `IOLocal`)

OpenTelemetry - Instrumentation API concepts

- `Context`: Immutable key-value pairs for storing all OTel-related objects
 - The `Context` object itself is propagated around (e.g. `ThreadLocal`, `IOLocal`)
 - e.g. `Span` object is stored in `Context`

OpenTelemetry - Instrumentation API concepts

- `Context`: Immutable key-value pairs for storing all OTel-related objects
 - The `Context` object itself is propagated around (e.g. `ThreadLocal`, `IOLocal`)
 - e.g. `Span` object is stored in `Context`

```
val context = Map(  
    SPAN_KEY -> Span(traceId, spanId, ...),  
    BAGGAGE_KEY -> Baggage(..)  
)
```

OpenTelemetry - Instrumentation API concepts

- `Context`: Immutable key-value pairs for storing all OTel-related objects
 - The `Context` object itself is propagated around (e.g. `ThreadLocal`, `IOLocal`)
 - e.g. `Span` object is stored in `Context`
 - ```
val context = Map(
 SPAN_KEY -> Span(traceId, spanId, ...),
 BAGGAGE_KEY -> Baggage(..)
)
```
- `OpenTelemetry/otel4s`: Central service for handling reporting spans and metrics

# OpenTelemetry - Instrumentation API concepts

- `Context`: Immutable key-value pairs for storing all OTel-related objects
  - The `Context` object itself is propagated around (e.g. `ThreadLocal`, `IOLocal`)
  - e.g. `Span` object is stored in `Context`

```
val context = Map(
 SPAN_KEY -> Span(traceId, spanId, ...),
 BAGGAGE_KEY -> Baggage(..)
)
```

- `OpenTelemetry/Otel4s`: Central service for handling reporting spans and metrics
- `Tracer`: Use to create spans. Created from `OpenTelemetry` instance


# Otel4s - Minimal Example

```
1 import cats.effect._
2 import cats.implicit.*
3 import org.typelevel.otel4s.oteljava.context.Context
4 import org.typelevel.otel4s.oteljava.OtelJava
5 import org.typelevel.otel4s.trace.Tracer
```

```
1 for
2 otel4s <- OtelJava.global[IO] // Initialize otel-java backend of Otel4s
3 given Tracer[IO] <- otel4s.tracerProvider.get("my_app")
4
5 _ <- Tracer[IO].span("root").surround(// Start a span,
ending an IO
6 for
7 _ <- Tracer[IO].span("child_1").surround(
8 IO.println("work work")
9)
10 _ <- Tracer[IO].span("child_2").surround(
11 Tracer[IO].span("grandchild_1").surround(
12 IO.println("more work done")
13)
14)
15 yield ()
16)
17 yield ()
```



Trace Start **March 3 2024, 20:17:05.270** | Duration **26.19ms** | Services **1** | Depth **3** | Total Spans **4**



Service & Operation ⌵ ➤ ⌵ ➤ || 19.39ms 21.84ms 24.29ms





# Integrating with Java libraries

- In Otel4s, we use IOLocal to propagate context, but Java libs use opentelemetry-java (`ThreadLocal`)

# Integrating with Java libraries

- In Otel4s, we use IOLocal to propagate context, but Java libs use opentelemetry-java (`ThreadLocal`)
- 💡 Extract `Context` from IOLocal and set it up as ThreadLocal before calling the java code

# Integrating with Java libraries

- In Otel4s, we use IOLocal to propagate context, but Java libs use opentelemetry-java (ThreadLocal)
- 💡 Extract Context from IOLocal and set it up as ThreadLocal before calling the java code

```
1 import cats.mtl.Local
2 import io.opentelemetry.context.Context as JContext // Context type from
Tel
3
4 def blockingWithContext[A](use: => A)(using local: Local[IO, Context]):
=
5 for
6 context <- local.ask
7 result <- IO.blocking {
8 val jContext: JContext = context.underlying
9 val scope = jContext.makeCurrent() // Set the ThreadLocal
10 try
11 use
12 finally
13 scope.close() // Unset the ThreadLocal
14 }
15 yield result
```

# Integrating with Java libraries - Example

# Integrating with Java libraries - Example

- From cats-effect, use Java 11's `HttpClient` to call another service

# Integrating with Java libraries - Example

- From cats-effect, use Java 11's `HttpClient` to call another service
- **Expectations:** Downstream service continues the trace



# Integrating with Java libraries - Example


- From cats-effect, use Java 11's `HttpClient` to call another service
- **Expectations:** Downstream service continues the trace

```
1 import
otellemetry.instrumentation.httpClient.JavaHttpClientTelemetry
2
3 for
4 otel4s <- OtelJava.global[IO]
5 given Tracer[IO] <- otel4s.tracerProvider.get("my_app")
6 given Local[IO, Context] = otel4s.localContext
7
8 httpClient = HttpClient.newBuilder().build()
9 instrumentedHttpClient = JavaHttpClientTelemetry
10 .builder(GlobalOpenTelemetry.get())
11 .build()
12 .newHttpClient(httpClient)
13 myService = MyService(instrumentedHttpClient)
14
15 _ <- myService.doWorkAndMakeRequest
16 yield ()
```

# Integrating with Java libraries - Example

```
1 class MyService(javaHttpClient: HttpClient)(using Local[IO, Context],
[IO]):
2
3 def doWorkAndMakeRequest: IO[Unit] =
4 withSpan("root")(for
5 _ <- withSpan("child_1")(IO.println("work work"))
6
7 _ <- withSpan("child_2")(for
8 req <- IO(HttpRequest.newBuilder()
9 .uri(new URI("http://localhost:8080/example"))
10 .GET()
11 .build())
12
13 resp <- blockingWithContext {
14 javaHttpClient.send(req, BodyHandlers.ofString)
15 }
16 _ <- IO.println(resp.body())
17 yield ())
18
19 yield ()
20
21 def withSpan[F[_], A](name: String)(using tracer: Tracer[F])(io: F[A]):
22
23 tracer.span(name).surround(io)
```

Trace Start **March 13 2024, 19:59:29.193** | Duration **812.6ms** | Services **2** | Depth **5** | Total Spans **6**



| Service & Operation           | 0µs | 203.15ms | 406.3ms | 609.45ms | 812.6 |
|-------------------------------|-----|----------|---------|----------|-------|
| ▼ my_app root                 |     |          |         |          |       |
| my_app child_1                |     |          |         |          |       |
| ▼ my_app child_2              |     |          |         |          |       |
| ▼ my_app GET                  |     |          |         |          |       |
| ▼ other_app Http Server - GET |     |          |         |          |       |
| other_app remoteAPI.co...     |     |          |         |          |       |



# Final thoughts

- Otel4s supports metrics too!

# Final thoughts

- Otel4s supports metrics too!
- use OTel-java's `Context.taskWrapping` to wrap threadpools

# Final thoughts

- Otel4s supports metrics too!
- use OTel-java's `Context.taskWrapping` to wrap threadpools
  - Many async libraries allow you to pass in your own threadpool (`Executor`)

# Final thoughts

- Otel4s supports metrics too!
- use OTel-java's `Context.taskWrapping` to wrap threadpools
  - Many async libraries allow you to pass in your own threadpool (`Executor`)
- Prefer Manual instrumentation? (Instead of using OTel Java agent)



**Thanks to..**

# Thanks to..

- otel4s and OpenTelemetry

# Thanks to..

- otel4s and OpenTelemetry
- <https://github.com/keuhdall/otel4s-grafana-example>

# Thanks to..

- otel4s and OpenTelemetry
- <https://github.com/keuhdall/otel4s-grafana-example>
- Scala 3